

Policy Mirror Descent with Reversed KL Projection

Ruitong Huang

Joint work with: Jincheng Mei, Chenjun Xiao, and Dale Schuurmans

Borealis AI

1 Background

- Reinforcement learning
- Policy gradient methods
 - REINFORCE
 - UREX


2 Our Algorithms

- PMD
- REPMD

3 Experimental Results

- Algorithmic tasks in OpenAI gym
- Ablation Study


Reinforcement Learning



- MDP: finite actions $A_t \in \mathcal{A}$ and finite states $S_t \in \mathcal{S}$;
- Agent: policy $\pi_{\theta}(a_t|s_t)$, parametrized by θ ; **Non-convex!**
- Environment: reward $\mathbb{P}(R_{t+1}|A_t, S_t)$, and transition $\mathbb{P}(S_{t+1}|A_t, S_t)$;

[Image source: Sutton and Barto, Reinforcement Learning book]

Reinforcement Learning



- WLOG, assume the state transition function is deterministic.
- State-action sequence: $\rho = (s_1, a_1, \dots, a_{T-1}, s_T)$;

Goal: Maximizing the expected reward

$$\max_{\theta} \mathbb{E} [R(\rho) \mid \pi_{\theta}] .$$

[Image source: Sutton and Barto, Reinforcement Learning book]

Goal: $\max_{\theta} \mathbb{E} [R(\rho) | \pi_{\theta}] = \max_{\theta} \sum_{\rho} \mathbb{P}(\rho | \pi_{\theta}) R(\rho);$

Compute the gradient: $\nabla f = f \nabla \log(f),$

$$\begin{aligned} \nabla_{\theta} \mathbb{E} [R(\rho) | \pi_{\theta}] &= \sum_{\rho} \nabla_{\theta} \mathbb{P}(\rho | \pi_{\theta}) R(\rho) = \sum_{\rho} \mathbb{P}(\rho | \pi_{\theta}) \nabla_{\theta} \log \mathbb{P}(\rho | \pi_{\theta}) R(\rho) \\ &= \mathbb{E}_{\rho} [\nabla_{\theta} \log \mathbb{P}(\rho | \pi_{\theta}) R(\rho)] \end{aligned}$$

To compute $\nabla_{\theta} \log \mathbb{P}(\rho | \pi_{\theta})$:

$$\nabla_{\theta} \log \mathbb{P}(\rho | \pi_{\theta}) = \nabla_{\theta} \sum_{t=1}^{T-1} \log \pi_{\theta}(a_t | s_t); \text{ (Dynamics Independent!)}$$

REINFORCE

REINFORCE

Compute the gradient w.r.t. θ :

- Sample ρ_1, \dots, ρ_M from the current π_θ ;
- Estimate the gradient $\mathbb{E}_\rho [\nabla_\theta \log \mathbb{P}(\rho | \pi_\theta) R(\rho)]$ by

$$\frac{1}{M} \sum_{\rho_i} R(\rho_i) \sum_{t=1}^{T-1} \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t});$$

- Update θ .

REINFORCE with entropy regularization

Entropy regularization to encourage exploration:

$$\max_{\theta} \mathbb{E}[R(\rho) | \pi_\theta] + \tau \mathcal{H}(\pi_\theta) = \max_{\theta} \mathbb{E}[R(\rho) | \pi_\theta] - \tau \pi_\theta \log \pi_\theta.$$

UREX

Reward-guided exploration:

$$\max_{\theta} \mathbb{E}[R(\rho) \mid \pi_{\theta}] + \tau \pi_{\tau} \log \pi_{\theta},$$

where $\pi_{\tau}(\rho) = \frac{1}{Z} \exp(\frac{1}{\tau} R(\rho))$.

- Encourages exploration, but emphasizes more on larger rewarded ρ ;
- May converge to poor local optima: estimate of π_{τ} is bad at the beginning of the training, thus it provides poor guidance;
- Difficult to interpret its fixed point (for a fixed τ), even in the simplex set;

$$\pi_{\theta}(\rho) = \frac{\tau \pi_{\tau}(\rho)}{\alpha - R(\rho)};^*$$

- No theoretical justification in its performance.

* α is the normalizer.

This work

- Typical analysis only focus on policy space.
- We focus on analysis in the parameter space.
 - One-layer neural network: $\pi_{\theta}(\cdot|s) = \text{softmax}(X_s^{\top} \theta)$;
 - Fixed points analysis;
 - Theoretical guarantee in performances.

Our Algorithms

Policy Mirror Descent (PMD)

PMD

Mirror descent type regularization (trust region):

$$\max_{\theta} \mathbb{E}[R(\rho) \mid \pi_{\theta}] - \tau D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\theta_{\text{old}}}).$$

- Essentially the same as TRPO;
- Our analysis is in the parameter space.

Projection view of PMD

$$\arg \min_{\pi_{\theta} \in \Pi} D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\tau}^*) \quad (\text{Projection})$$

$$\text{where } \pi_{\tau}^* = \arg \max_{\pi \in \Delta} \mathbb{E}[R(\rho) \mid \pi] - \tau D_{\text{KL}}(\pi \parallel \pi_{\theta_{\text{old}}}) \quad (\text{Improvement})$$

Properties of PMD

If the projection step can be solved optimally,

- Property 1: Monotonically improves the expected reward;

$$\mathbb{E}[R(\rho) \mid \pi_{\theta_{t+1}}] - \mathbb{E}[R(\rho) \mid \pi_{\theta_t}] \geq 0,$$

- Property 2: Its fixed points include the optimal policy in the parameter space.
-
- TRPO also has the monotonical improvement property, but only in the policy space;
 - Poor local optima still exist: can be observed in simulation and real experiments;
 - In practice, the projection step cannot be solved optimally, because it is non-convex in θ , even for 1-layer neural network.

Reversed Entropy Policy Mirror Descent (REPMD)

REPMD

$$\arg \min_{\pi_{\theta} \in \Pi} D_{\text{KL}}(\pi_{\tau, \tau'}^* || \pi_{\theta})$$

$$\text{where } \pi_{\tau, \tau'}^* = \arg \max_{\pi \in \Delta} \mathbb{E}[R(\rho) | \pi_{\theta}] - \tau D_{\text{KL}}(\pi_{\theta} || \pi_{\theta_{\text{old}}}) + \tau' \mathcal{H}(\pi_{\theta}),$$

with $\tau \rightarrow \infty$ and $\tau' \rightarrow 0$.

- **Entropy regularization:** encourages exploration to mitigate the local optima problem;
- **Reversed entropy:** convexifies the projection step in θ for the 1-layer neural network case, thus the projection step is solvable even in the parameter space.

Properties of REPMD

When using 1-layer neural network: $\pi_{\theta}(\cdot|s) = \text{softmax}(X_s^{\top} \theta)$, REPMD satisfies

- Property 1: Fixed points include the optimal policy.
- Property 2: Monotonically improves the surrogate reward;

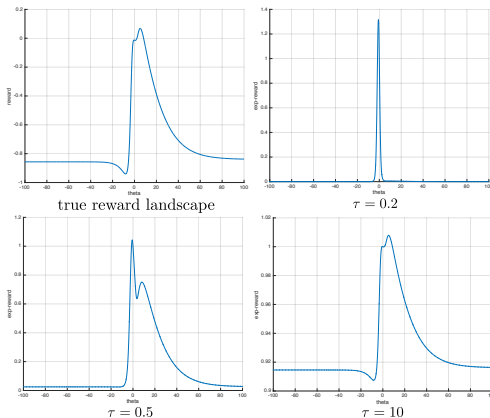
$$SR(\pi_{\theta_{t+1}}) - SR(\pi_{\theta_t}) \geq 0,$$

$$\text{where } SR(\pi) = (\tau + \tau') \log \sum_{\rho} \exp \left(\frac{R(\rho) + \tau \log \pi(\rho)}{\tau + \tau'} \right).$$

- Local optima still exist.

Behavior of $SR(\pi)$

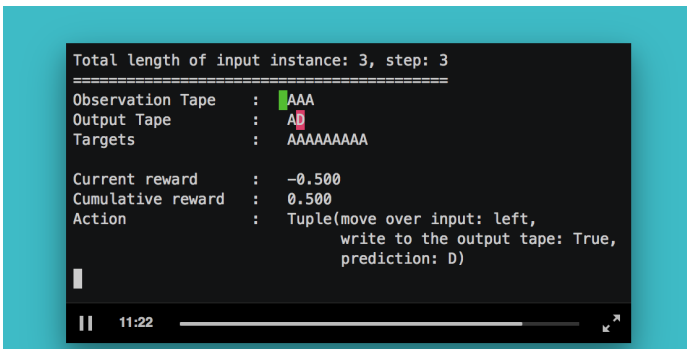
- $SR(\pi) = \tau \log \sum_{\rho} \pi(\rho) \exp \left(\frac{R(\rho)}{\tau} \right)$, as $\tau' \rightarrow 0$;
- $SR(\pi) \rightarrow \mathbb{E}[R(\rho)|\pi]$, as $\tau \rightarrow \infty$ and $\tau' \rightarrow 0$.
- Landscape of $SR(\pi)$: X-axis is θ , Y-axis is reward.



Experiments

- Tasks:
 - Bandit simulation
 - OpenAI gym algorithmic tasks
 - **Copy:** $ababaa \rightarrow ababaa$.
 - **DuplicatedInput:** $aba \rightarrow aabbaa$.
 - **RepeatCopy:** $abc \rightarrow abccbaabc$.
 - **Reverse:** $aabc \rightarrow cbaa$.
 - **ReversedAddition:** Observe two numbers in base 3 in little-endian order on a $2 \times n$ grid tape: $120; 221 \rightarrow 022$
- Reward:
 - Correct(+1); incorrect(-0.5 & terminate); idle(0); overtime(-1);
 - Only accumulative reward is available
- Observation: letter on the current position of the read head
- Actions: move the read head; write (some symbol or nothing);

Experiments

A screenshot of a terminal window with a black background and white text, set against a teal background. The terminal displays the state of an OpenAI Gym environment. At the top, it says 'Total length of input instance: 3, step: 3' followed by a line of equals signs. Below this, it lists 'Observation Tape : AAA' with a green cursor at the first 'A', 'Output Tape : AD' with a red cursor at the 'D', and 'Targets : AAAAAAAAAA'. Further down, it shows 'Current reward : -0.500', 'Cumulative reward : 0.500', and 'Action : Tuple(move over input: left, write to the output tape: True, prediction: D)'. At the bottom left, there is a small white bar and a pause icon. At the bottom right, there is a timestamp '11:22' and a progress bar.

```
Total length of input instance: 3, step: 3
=====
Observation Tape : AAA
Output Tape      : AD
Targets          : AAAAAAAAAA

Current reward    : -0.500
Cumulative reward : 0.500
Action            : Tuple(move over input: left,
                           write to the output tape: True,
                           prediction: D)

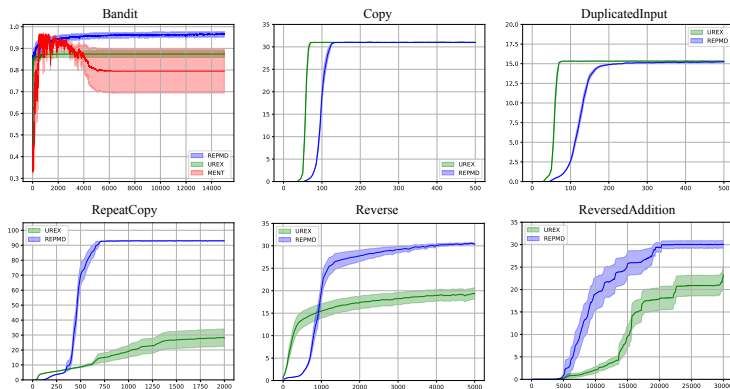
[Bar]

|| 11:22 [Progress Bar]
```


- Reward:
 - Correct(+1); incorrect(-0.5 & terminate); idle(0); overtime(-1);
 - Only accumulative reward is available
- Observation: letter on the current position of the read head
- Actions: move the read head; write (some symbol or nothing);

[Image source: OpenAI Gym]

Experiments



- X-axis is number of sampled trajectories; Y-axis is reward;
- REINFORCE (red), UREX (green), and REPMD (blue);
- Bandit results averaged over 5 repetitions. Algorithmic task results averaged over 25 runs (5 repetitions \times 5 random seeds).



- REPMD (blue), REPMD w/o entropy (yellow), PMD with entropy (red), REPMD w/o projection (green).

Thank you !
Questions?